

Enrolment No./Seat No\_\_\_\_\_

## GUJARAT TECHNOLOGICAL UNIVERSITY

BE - SEMESTER-III (NEW) EXAMINATION – SUMMER 2024

**Subject Code:3130703**

**Date:29-06-2024**

**Subject Name: Database Management Systems**

**Time:10:30 AM TO 01:00 PM**

**Total Marks:70**

**Instructions:**

1. Attempt all questions.
2. Make suitable assumptions wherever necessary.
3. Figures to the right indicate full marks.
4. Simple and non-programmable scientific calculators are allowed.

|            |  | MARKS     |
|------------|--|-----------|
| <b>Q.1</b> | (a) Define Logical Data Independence and Physical Data Independence.   | <b>03</b> |
|            | (b) Draw and explain 3 layers of abstraction of DBMS   | <b>04</b> |
|            | (c) Write relational algebra expression for following<br><b>(query 1,2 and 3 carry one mark/query 4 &amp; 5 carry two marks)</b><br>Employee(eid, ename, salary, deptid)<br>Department(deptid, dname)<br>1) List name of all employee<br>2) List all the departments with id > 5<br>3) List sum of the salary all employees.<br>4) Department wise count the number of employees<br>5) Find the name of employee earning highest salary. | <b>07</b> |
| <b>Q.2</b> | (a) List various symbol used for ER diagram  | <b>03</b> |
|            | (b) Define referential integrity constraint with example.  | <b>04</b> |
|            | (c) List the Armstrong axioms used to compute closure of the set of functional dependency.   | <b>07</b> |
|            | <b>OR</b>  |           |
|            | (c) What is decomposition? Explain dependency preserving decomposition with the help of example.   | <b>07</b> |
| <b>Q.3</b> | (a) Give syntax of command used to add a new column in a table.  | <b>03</b> |
|            | (b) Define indexing. How will it be useful for searching data?   | <b>04</b> |
|            | (c) How can we measure cost of Join query? Explain using example.  | <b>07</b> |
|            | <b>OR</b>  |           |
| <b>Q.3</b> | (a) Write an SQL command to modify the content of table.   | <b>03</b> |
|            | (b) Explain sparse indexing with the help of diagram.  | <b>04</b> |
|            | (c) Draw and explain the query processing steps  | <b>07</b> |
| <b>Q.4</b> | (a) What is relational model? Explain schema diagram using example.  | <b>03</b> |
|            | (b) Define transaction and its properties?   | <b>04</b> |
|            | (c) Write a note on view serializability   | <b>07</b> |
|            | <b>OR</b>  |           |
| <b>Q.4</b> | (a) With the help of example explain Generalization..  | <b>03</b> |
|            | (b) Define various approach used for Log based recovery  | <b>04</b> |
|            | (c) List and explain various types of Locks used for concurrency control.  | <b>07</b> |

- |            |   |           |
|------------|---|-----------|
| <b>Q.5</b> | <b>(a)</b> How can we print output in PL/SQL block? Give example. | <b>03</b> |
|            | <b>(b)</b> Explain cursor and its types.                          | <b>04</b> |
|            | <b>(c)</b> Write SQL queries for the following                    | <b>07</b> |

**(query 1,2 and 3 carry one mark/query 4 & 5 carry two marks)**

Employee(eid, ename, salary, deptid)

Department(deptid, dname)

1. List name of all employee
2. List all the departments with id > 5
3. List sum of the salary all employees.
4. Department wise count the number of employees
5. Find the name of employee earning highest salary.

**OR**

- |            |  |           |
|------------|--|-----------|
| <b>Q.5</b> | <b>(a)</b> Define encryption and decryption.   | <b>03</b> |
|            | <b>(b)</b> List various types of Triggers.   | <b>04</b> |
|            | <b>(c)</b> What is role? How can we authorize a user using grant and revoke command? Give example. | <b>07</b> |

\*\*\*\*\*

**Q.1 (a) Define Logical Data Independence and Physical Data Independence. (3 mark)**

### **Introduction / Definition:**

In database management systems (DBMS), **data independence** refers to the ability to **change the schema at one level without affecting the schema at the next higher level**. It is a key feature of DBMS that allows flexibility and reduces maintenance efforts. There are two types of data independence: **Logical Data Independence** and **Physical Data Independence**.

---

#### **1. Logical Data Independence:**

- Refers to the **ability to change the logical schema** (table structure, attributes, relationships) **without affecting the application programs** or external schema.
  - Example: Adding a new column to a table or creating a new table does not require changes in the application using existing data.
  - **Purpose:** Protects applications from frequent changes in **logical design**.
- 

#### **2. Physical Data Independence:**

- Refers to the **ability to change the physical storage structure** (file organization, indexing, storage devices) **without affecting the logical schema or application programs**.
- Example: Changing from sequential file storage to indexed storage does not impact how the data is viewed or queried.
- **Purpose:** Ensures flexibility in **storage optimization** and improves performance without affecting users.

---

## **Key Points:**

- Logical independence deals with **changes in table structures and relationships**.
- Physical independence deals with **changes in storage and access methods**.
- Both types are crucial for **maintaining database consistency, flexibility, and ease of maintenance**.

(b) Draw and explain 3 layers of abstraction of DBMS (4 mark)

### **Introduction / Definition:**

Database Management Systems (DBMS) use **abstraction** to simplify interaction with data. Abstraction hides the complexities of data storage and provides **different views for different users**. There are **three levels of abstraction**: Physical Level, Logical Level, and View Level. These layers ensure **data independence, security, and ease of access**.

---

### **1. Physical Level (Lowest Level)**

- Describes **how data is actually stored** in the database (files, indexes, storage paths).
- Focuses on **physical storage structures and access methods**.
- Provides details like **record placement, data blocks, and pointers**.
- Users **do not see this layer**; it is managed by the DBMS.

---

### **2. Logical Level (Conceptual Level / Middle Level)**

- Describes **what data is stored and relationships among data**.

- Hides **physical storage details** from users.
  - Defines tables, attributes, data types, and constraints.
  - Provides a **community-wide view** for database designers and administrators.
  - Example: Customer table has attributes Cust\_Id, Name, Age, Address, Salary.
- 

### 3. View Level (Highest Level / External Level)

- Describes **how data is viewed by individual users or applications.**
  - Each user can have **customized views** based on their needs.
  - Hides unnecessary details and **presents only relevant data.**
  - Example: HR department may see employee name and salary, but not personal address.
- 

#### Key Points / Advantages:

- **Data Abstraction** simplifies database usage.
  - Provides **data security** by restricting access at the view level.
  - Supports **logical and physical data independence.**
  - Allows multiple users to interact with the database **without interfering with each other.**
- 

#### Diagram (Described in Words):

- **Top Layer (View Level):** User-specific views → Shows only required data.

- **Middle Layer (Logical Level):** Conceptual schema → Defines tables, attributes, and relationships.
  - **Bottom Layer (Physical Level):** Physical storage → Files, indexes, and storage details.  
*(Imagine a three-layered pyramid with View at the top, Logical in the middle, and Physical at the bottom.)*
- 

(c) Write relational algebra expression for following (query 1,2 and 3 carry one mark/query 4 & 5 carry two marks) Employee(eid, ename, salary, deptid) Department(deptid, dname)

- 1) List name of all employee
- 2) List all the departments with id > 5
- 3) List sum of the salary all employees.
- 4) Department wise count the number of employees
- 5) Find the name of employee earning highest salary.

## Relational Algebra Expressions

Given Tables:

- **Employee(eid, ename, salary, deptid)**
  - **Department(deptid, dname)**
- 

### 1) List name of all employees (1 mark)

$$\pi_{ename}(\text{Employee})$$

- **Explanation:** The  $\pi$  (**projection**) operator selects the ename column from the Employee table.

### 2) List all departments with deptid > 5 (1 mark)

$$\sigma_{deptid>5}(\text{Department})$$

- **Explanation:** The  **$\sigma$  (selection)** operator retrieves rows where deptid > 5.
- 

### 3) List sum of the salary of all employees (1 mark)

$$\gamma_{\text{SUM}(\text{salary})}(\text{Employee})$$

- **Explanation:** The  **$\gamma$  (aggregate)** operator calculates the sum of the salary attribute for all employees.
- 

### 4) Department-wise count of number of employees (2 marks)

$$\gamma_{\text{deptid}, \text{COUNT}(\text{eid})}(\text{Employee})$$

- **Explanation:**
    - $\gamma$  is the **grouping operator**.
    - Groups employees by deptid and counts eid in each department.
    - Returns **deptid** and number of employees in each department.
- 

### 5) Find the name of employee earning highest salary (2 marks)

$$\pi_{ename}(\sigma_{\text{salary}=\text{MAX}(\text{salary})}(\text{Employee}))$$

- **Explanation:**
  - $\sigma$  selects rows where salary equals the maximum salary.

- $\pi$  projects only the ename column.
- This retrieves the name(s) of employee(s) with the highest salary.

**Q.2 (a) List various symbol used for ER diagram (3 mark)**

### **Introduction / Definition:**

An **ER diagram** is a graphical representation of entities, their attributes, and relationships in a database. It is widely used during **database design** to visually depict how data is organized and interrelated. ER diagrams use **specific symbols** to represent entities, relationships, and attributes.

---

### **Common Symbols in ER Diagram:**

#### **1. Rectangle (Entity)**

- Represents an **entity**, which is a real-world object or concept.
- Example: Employee, Department.

#### **2. Oval / Ellipse (Attribute)**

- Represents an **attribute** of an entity or relationship.
- Example: Emp\_Id, Name, Salary.

#### **3. Diamond (Relationship)**

- Represents a **relationship** between two or more entities.
- Example: Works\_In between Employee and Department.

#### **4. Double Oval (Multivalued Attribute)**

- Represents an attribute that can have **multiple values**.
- Example: PhoneNumbers for an employee.

#### **5. Double Rectangle (Weak Entity)**

- Represents a **weak entity**, which depends on a strong entity for its identification.
- Example: Dependent of an Employee.

## 6. Dashed Oval (Derived Attribute)

- Represents an attribute that can be **derived** from other attributes.
- Example: Age can be derived from DateOfBirth.

## 7. Lines / Connectors

- Connect entities to attributes or relationships.
- Shows participation in relationships.

## 8. Underlined Attribute

- Denotes the **primary key** of an entity.
- Example: Emp\_Id underlined in Employee.

### Key Points:

- These symbols provide a **standardized way** to represent data structures visually.
- Help in understanding **data relationships and constraints** before actual database implementation.
- Make **database design clear, easy to communicate, and error-free**.

(b) Define referential integrity constraint with example. (4 mark)

### Introduction / Definition:

In a relational database, **referential integrity** is a type of **integrity constraint** that ensures the **consistency and correctness of data across related tables**. It enforces that a **foreign key in one table**

must either **match a primary key in another table** or be **NULL**. This prevents invalid data and maintains relationships between tables.

---

### **Explanation:**

- A **foreign key** in a child table refers to the **primary key** in a parent table.
  - Referential integrity ensures that **every value of the foreign key exists in the parent table**.
  - Violating this rule can lead to **orphan records**, which are child records without a corresponding parent record.
  - SQL provides options such as **ON DELETE CASCADE** or **ON UPDATE CASCADE** to handle changes in the parent table.
- 

### **Example:**

#### **Tables:**

- **Department(deptid, dname)** → deptid is primary key
- **Employee(eid, ename, deptid)** → deptid is foreign key referencing Department

#### **SQL Constraint:**

```
CREATE TABLE Department (
    deptid NUMBER PRIMARY KEY,
    dname VARCHAR2(50)
);
```

```
CREATE TABLE Employee (
```

```
eid NUMBER PRIMARY KEY,  
ename VARCHAR2(50),  
deptid NUMBER,  
CONSTRAINT fk_dept FOREIGN KEY (deptid)  
REFERENCES Department(deptid)  
);
```

- Here, every deptid in Employee **must exist** in the Department table.
- If someone tries to insert an employee with deptid = 10 but no such department exists, the DBMS will **reject the insert**.

---

### Key Points / Advantages:

- Ensures **data consistency** across tables.
- Prevents **orphan records** in child tables.
- Maintains **correct relationships** in relational databases.

(c) List the Armstrong axioms used to compute closure of the set of functional dependency. (7 mark)

### Introduction / Definition:

In a relational database, **functional dependency (FD)** describes a relationship between attributes where the value of one set of attributes determines the value of another. To **compute the closure of a set of functional dependencies** or to infer all possible FDs from a given set, **Armstrong's axioms** are used. These axioms are **sound and complete rules** for reasoning about functional dependencies.

---

### Armstrong's Axioms

There are **three primary axioms** and **three additional rules derived from them**:

---

### 1. Reflexivity Rule:

- If  $Y$  is a subset of  $X$ , then  $X \rightarrow Y$
  - **Explanation:** Every set of attributes functionally determines its subset.
  - **Example:** If  $X = \{A, B\}$ , then  $\{A, B\} \rightarrow \{A\}$ .
- 

### 2. Augmentation Rule:

- If  $X \rightarrow Y$ , then  $XZ \rightarrow YZ$  for any  $Z$
  - **Explanation:** Adding the same attributes to both sides of a dependency preserves the dependency.
  - **Example:** If  $A \rightarrow B$ , then  $AC \rightarrow BC$ .
- 

### 3. Transitivity Rule:

- If  $X \rightarrow Y$  and  $Y \rightarrow Z$ , then  $X \rightarrow Z$
  - **Explanation:** Functional dependency is transitive.
  - **Example:** If  $A \rightarrow B$  and  $B \rightarrow C$ , then  $A \rightarrow C$ .
- 

## Additional Rules (Derived from the above)

### 4. Union Rule:

- If  $X \rightarrow Y$  and  $X \rightarrow Z$ , then  $X \rightarrow YZ$
- **Explanation:** Two dependencies with the same determinant can be combined.

## **5. Decomposition Rule:**

- If  $X \rightarrow YZ$ , then  $X \rightarrow Y$  and  $X \rightarrow Z$
- **Explanation:** A dependency on multiple attributes can be broken down into multiple dependencies.

## **6. Pseudotransitivity Rule:**

- If  $X \rightarrow Y$  and  $WY \rightarrow Z$ , then  $XW \rightarrow Z$
  - **Explanation:** Combines transitivity and augmentation in a single step.
- 

### **Key Points / Features:**

- **Soundness:** All derived FDs are valid.
  - **Completeness:** All valid FDs can be derived using these axioms.
  - Used for **computing closure, normalization, and schema design.**
- 

### **Example of Computing Closure:**

- Given FDs:  $\{A \rightarrow B, B \rightarrow C\}$
- Using **Transitivity**,  $A \rightarrow C$
- Using **Union**,  $A \rightarrow BC$

**OR**

(c) What is decomposition? Explain dependency preserving decomposition with the help of example. (7 mark)

### **Introduction / Definition:**

In database design, **decomposition** refers to the process of **breaking a relation into two or more smaller relations** while preserving

important properties such as **lossless join** and **dependency preservation**. Decomposition is commonly used during **normalization** to eliminate redundancy, avoid anomalies, and maintain **data integrity**.

---

## Dependency-Preserving Decomposition

### Definition:

A decomposition is said to be **dependency-preserving** if all **functional dependencies** (FDs) from the original relation can be enforced by applying FDs **only on the smaller relations** after decomposition, without needing to join them back together.

### Importance:

- Ensures **constraints and rules** on data are still enforceable after decomposition.
  - Reduces the need for expensive **joins** to check functional dependencies.
  - Helps maintain **data consistency and integrity**.
- 

### Steps / Procedure:

1. Identify all **functional dependencies (FDs)** in the original relation.
2. Check if the relation violates **normal forms** (1NF, 2NF, 3NF, BCNF).
3. Decompose the relation into smaller relations to **remove redundancy** and satisfy the desired normal form.
4. Ensure the decomposition is **lossless** (can reconstruct original relation using joins).

5. Check if the decomposition is **dependency-preserving**. If yes, all original FDs can be enforced on the new relations.
- 

### **Example:**

#### **Original Relation:**

$R(A, B, C)$  with FDs:

- $A \rightarrow B$
- $B \rightarrow C$

#### **Step 1: Identify FDs violating BCNF**

- $B \rightarrow C$  violates BCNF because B is not a superkey.

#### **Step 2: Decompose Relation**

- Decompose R into:
  1.  $R1(B, C)$
  2.  $R2(A, B)$

#### **Step 3: Check Dependency Preservation**

- Original FDs:  $A \rightarrow B$  and  $B \rightarrow C$
- $A \rightarrow B$  is preserved in  $R2$
- $B \rightarrow C$  is preserved in  $R1$
- All FDs are preserved without joining  $R1$  and  $R2$

#### **Result:**

- Decomposition is **lossless** and **dependency-preserving**.
- 

#### **Key Points / Advantages:**

- Reduces **redundancy and anomalies** (insertion, update, deletion).
- Ensures **data integrity** through preserved dependencies.
- Makes database **efficient and consistent**.

### **Disadvantages:**

- Sometimes achieving both **lossless and dependency-preserving** decomposition may not be possible for all relations.
  - May require **more joins** in queries if relations are decomposed excessively.
- 

**Q.3 (a) Give syntax of command used to add a new column in a table.  
(3 mark)**

### **Adding a New Column in a Table**

#### **Introduction / Definition:**

In SQL, the structure of a table can be **modified after creation** using the ALTER TABLE command. One common operation is to **add a new column** to an existing table. This allows the database to **accommodate new data requirements** without dropping or recreating the table.

---

#### **Syntax:**

```
ALTER TABLE table_name
ADD column_name datatype [constraints];
```

#### **Explanation of Syntax:**

- ALTER TABLE table\_name → Specifies the table you want to modify.

- ADD → Keyword used to add a new column.
  - column\_name → Name of the new column to be added.
  - datatype → Specifies the type of data the column will store (e.g., NUMBER, VARCHAR2, DATE).
  - [constraints] → Optional; can include NOT NULL, UNIQUE, DEFAULT, etc.
- 

### **Example:**

Suppose we have a table Employee(eid, ename, salary) and we want to add a deptid column:

```
ALTER TABLE Employee
```

```
ADD deptid NUMBER;
```

### **Explanation:**

- Adds a new column deptid of type NUMBER to the Employee table.
  - Existing rows will have NULL in the new column until updated.
- 

### **Key Points / Advantages:**

- Allows **flexible modification** of table structure.
- Avoids the need to **drop and recreate** tables for small changes.
- Can add **multiple columns** in a single ALTER TABLE statement if needed.

**(b) Define indexing. How will it be useful for searching data? (4 mark)**

### **Introduction / Definition:**

**Indexing** is a database technique used to **speed up data retrieval**

from a table. An index is a **data structure** (like a pointer or a table of keys) that allows the DBMS to **locate records quickly** without scanning the entire table. It is similar to an index in a book, which helps find information without reading every page.

---

### **How Indexing Works:**

- When a table is indexed on a column, the DBMS creates a **separate data structure** storing the column values along with **pointers to the corresponding rows** in the table.
  - During a query, instead of scanning all rows, the DBMS searches the **index structure** to directly find the required data.
  - Common types of indexes include **single-column index**, **composite index**, **unique index**, and **clustered/non-clustered index**.
- 

### **Uses / Advantages for Searching Data:**

- **Faster Search:** Reduces the time to locate rows, especially in **large tables**.
  - **Efficient Query Execution:** Improves performance of SELECT, WHERE, and JOIN queries.
  - **Supports Sorting:** Helps in **ORDER BY** and **GROUP BY** operations.
  - **Reduces Disk I/O:** DBMS reads fewer blocks, saving memory and processing time.
- 

### **Example:**

Suppose we have a table Employee(eid, ename, salary) and we frequently search by eid.

```
CREATE INDEX idx_eid ON Employee(eid);
```

- Now, a query like `SELECT * FROM Employee WHERE eid = 101;` will **use the index** to quickly locate the row instead of scanning the entire table.
- 

## Key Points:

- Indexing improves **query performance** but may **slightly slow down insert, update, and delete operations** because the index must also be updated.
- Proper use of indexes is critical for **database optimization**.

(c) How can we measure cost of Join query? Explain using example.

(7 mark)

## Introduction / Definition:

In relational databases, **join operations** combine rows from two or more tables based on a related column. Join queries are fundamental but can be **resource-intensive**, especially for large tables. Measuring the **cost of a join query** is essential for **query optimization**. The cost generally refers to the **amount of I/O, CPU, and memory resources** needed to execute the join.

---

## Factors Affecting Join Cost:

- Number of tuples (rows)** in each table.
- Size of each tuple** (number of attributes and data types).
- Access method** (table scan, index scan).
- Join algorithm** used (nested loop, sort-merge, hash join).

## 5. Available memory and buffer pages for temporary storage.

---

### Join Algorithms and Cost Estimation:

#### 1. Nested Loop Join (NLJ):

- **Method:** For each row in the outer table, scan the inner table to find matching rows.
- **Cost Formula:**

$$\begin{aligned} \text{Cost} \\ = & (\text{Number of rows } o) \times (\text{Cost of access } i) \end{aligned}$$

$$\begin{aligned} & f r \\ & o a \end{aligned}$$

ows i  
ccess i

- **Example:**
  - Table A: 1000 rows, Table B: 500 rows
  - If scanning B fully for each row in A, cost =  $1000 \times 500 = 500,000$  row accesses.

---

#### 2. Sort-Merge Join (SMJ):

- **Method:** Sort both tables on the join key, then merge them to find matching tuples.
- **Cost Formula:**

$$\begin{aligned} \text{Cost} \\ = & \text{Cost of sorting } o \\ + & \text{Cost of merging } o \end{aligned}$$

$$\begin{aligned} & f s \\ & f m \end{aligned}$$

orting  
ergin

- **Example:**
  - Table A: 1000 rows, Table B: 500 rows

- Sort A and B, then merge → fewer comparisons than nested loop, especially for large tables.
- 

### 3. Hash Join:

- **Method:** Build a hash table on the smaller table using the join key, then probe it with rows from the larger table.
  - **Cost Formula:**
  - Cost=Cost to build hash table+Cost to probe hash table
  - **Example:**
    - Table A: 1000 rows (build hash), Table B: 500 rows (probe)
    - Efficient for **equality joins** and large datasets.
- 

### Steps to Measure Join Cost:

1. Determine **join algorithm** based on table sizes and indexes.
  2. Estimate **number of I/O operations** needed to read tables.
  3. Add **CPU cost** for comparisons and hashing/sorting.
  4. Compute **total cost** = I/O cost + CPU cost + memory usage.
  5. Compare costs for different algorithms to choose **optimal join plan**.
- 

### Example in Words:

Suppose we have two tables:

- **Employee(eid, ename, deptid)** → 1000 rows
- **Department(deptid, dname)** → 50 rows

Query:

```
SELECT e.ename, d.dname  
FROM Employee e  
JOIN Department d  
ON e.deptid = d.deptid;
```

### Cost Estimation:

- Using **nested loop join** with Department as outer:
  - Cost =  $50 \times 1000 = 50,000$  row accesses
- Using **hash join** with Department as build table:
  - Cost = build hash on 50 rows + probe 1000 rows → much lower than nested loop.

---

### Key Points / Advantages of Measuring Join Cost:

- Helps **query optimizer** choose the most efficient execution plan.
- Reduces **query execution time** and resource usage.
- Essential for **large databases** and multi-user systems.

OR

**Q.3 (a) Write an SQL command to modify the content of table. (3 mark)**

### Introduction / Definition:

In SQL, the **content of a table** can be modified using the UPDATE command. This command allows changing **existing data** in one or more rows of a table based on a specified condition. It is an essential part of **data manipulation operations (DML)** in a relational database.

---

## Syntax:

```
UPDATE table_name  
SET column1 = value1, column2 = value2, ...  
WHERE condition;
```

## Explanation of Syntax:

- UPDATE table\_name → Specifies the table whose data needs to be modified.
  - SET column1 = value1, ... → Lists the columns to be updated and the new values.
  - WHERE condition → Optional; specifies which rows to update. If omitted, **all rows** will be updated.
- 

## Example:

Suppose we have a table Employee(eid, ename, salary, deptid) and we want to **increase the salary of employee with eid = 101**:

```
UPDATE Employee  
SET salary = 75000  
WHERE eid = 101;
```

## Explanation:

- Updates the salary column for the employee whose eid is 101.
  - Other rows in the table remain unaffected.
- 

## Key Points / Advantages:

- Allows **selective updating** of records.

- Supports **updating multiple columns** at once.
- Helps maintain **current and accurate data** in the database.

(b) Explain sparse indexing with the help of diagram. (4 mark)

**Indexing** in databases is a technique to speed up data retrieval. A **sparse index** is a type of **non-dense indexing** in which **index entries are not created for every record in the data file**, but only for **some records**, usually one per data block. Sparse indexing reduces the **size of the index** and **speeds up search** while using less storage.

---

### **Explanation:**

- In a sparse index, the index stores **pointers to only some records** (typically the first record of each block).
  - When searching, the DBMS uses the sparse index to **locate the block** and then scans that block to find the exact record.
  - Sparse indexing is efficient for **large files** where creating a dense index for every record would be costly.
  - Typically used with **sequentially ordered data**.
- 

### **Advantages:**

- Requires **less storage** than dense indexing.
- Faster **index creation and maintenance**.
- Reduces **disk I/O** compared to scanning the entire table.

### **Disadvantages:**

- Searching may require **scanning within a block** after finding the block via the index.
- Slightly slower than dense indexing for exact record retrieval.

---

### **Example / Diagram (Described in Words):**

Consider a data file with records sorted by Emp\_Id:

#### **Data File (Blocks):**

- Block 1 → Emp\_Id: 101, 102, 103
- Block 2 → Emp\_Id: 104, 105, 106
- Block 3 → Emp\_Id: 107, 108, 109

#### **Sparse Index:**

- Index entry → Points to the **first record in each block**:
  - 101 → Block 1
  - 104 → Block 2
  - 107 → Block 3

#### **Search Process:**

1. To find Emp\_Id 105 → Check sparse index → 104 points to Block 2
2. Scan Block 2 → Locate 105

*(Imagine a table of Index → Block pointer mapping on the left and data blocks on the right.)*

**(c) Draw and explain the query processing steps (7 mark)**

#### **Introduction / Definition:**

**Query processing** in a Database Management System (DBMS) is the process of **translating a high-level SQL query into an efficient execution plan** that retrieves the desired data from the database. It ensures **correctness, efficiency, and optimization** of queries. Query processing involves several systematic steps to minimize **CPU, I/O, and memory costs**.

---

## Steps of Query Processing:

### 1. Parsing and Translation:

- The SQL query is **parsed** to check **syntax and validity**.
- **Semantic checks** ensure tables, attributes, and data types exist.
- The query is translated into an **internal representation**, typically **relational algebra**.
- **Example:** SELECT ename FROM Employee WHERE deptid=10;  
→ Translated to relational algebra:

$$\pi_{ename}(\sigma_{deptid=10}(Employee))$$

---

### 2. Query Optimization:

- The internal representation is passed to the **query optimizer**.
- The optimizer generates **equivalent query plans** using **algebraic transformations**.
- Selects the **most efficient execution plan** based on:
  - Available indexes
  - Join order
  - Estimated cost (I/O and CPU)

---

### 3. Plan Generation / Evaluation:

- The optimized query plan is converted into a **sequence of low-level operations** that the DBMS can execute.
- Operations include **table scans, index scans, joins, projections, and selections**.

- The DBMS may also determine **temporary storage** requirements and **access paths**.
- 

#### 4. Query Execution:

- The **execution engine** carries out the plan step by step.
  - Data is retrieved, joined, filtered, and projected as per the plan.
  - Final result is returned to the user or application.
- 

#### Flow Diagram (Described in Words):

1. **User Query (SQL)** →
2. **Parsing & Translation** → Check syntax and convert to relational algebra →
3. **Query Optimization** → Generate multiple plans → Choose optimal plan →
4. **Plan Generation** → Convert plan to executable operations →
5. **Query Execution** → Retrieve data → Return result

*(Imagine a flow from top to bottom with arrows connecting the above steps, showing SQL input to final output.)*

---

#### Key Points / Advantages:

- Ensures **efficient query execution** and **minimal resource usage**.
- Helps in **cost estimation** and selecting the **best join/order strategy**.
- Supports **large databases** with multiple users and complex queries.

---

## **Example in Words:**

- SQL Query: SELECT ename, salary FROM Employee WHERE salary > 50000;
- **Parsing:** Checks table Employee and columns exist.
- **Optimization:** Chooses to use an index on salary (if available) instead of full scan.
- **Execution:** Retrieves only rows with salary > 50000 and projects ename and salary.

**Q.4 (a) What is relational model? Explain schema diagram using example. (3 mark)**

### **Introduction / Definition:**

The **relational model** is a way to organize data in a database using **tables (relations)**, where each table consists of **rows (tuples)** and **columns (attributes)**. It was proposed by **E.F. Codd** in 1970 and forms the foundation of most modern database systems. The relational model ensures **data integrity, flexibility, and ease of querying** using SQL.

---

### **Key Features:**

- Data is stored in **relations (tables)**.
- Each row (tuple) represents a **unique record**.
- Each column (attribute) represents a **data field**.
- Relationships between tables are established using **primary and foreign keys**.
- Supports **operations like select, project, join** using relational algebra.

---

## **Relational Schema and Schema Diagram:**

- A **relational schema** defines the **structure of a relation**: its name, attributes, and constraints.
- **Schema Diagram** is a **graphical representation** of relations and their relationships.

### **Example:**

#### **1. Tables:**

- Employee(Emp\_Id, Ename, Dept\_Id, Salary)
- Department(Dept\_Id, Dname)

#### **2. Relationships:**

- Dept\_Id in Employee is a **foreign key** referencing Department.

## **Schema Diagram (Described in Words):**

- Rectangle → Employee with attributes: Emp\_Id (PK), Ename, Dept\_Id, Salary
- Rectangle → Department with attributes: Dept\_Id (PK), Dname
- Line connecting Employee.Dept\_Id → Department.Dept\_Id indicating **foreign key relationship**
- **Interpretation:**
  - Each employee belongs to one department.
  - One department can have multiple employees.

---

## **Key Points / Advantages:**

- Makes data **organized and structured**.

- Ensures **referential integrity** through primary and foreign keys.
- Easy to query using SQL.

(b) Define transaction and its properties? (4 mark)

### **Introduction / Definition:**

A **transaction** in a Database Management System (DBMS) is a **logical unit of work** that consists of one or more database operations such as **INSERT, UPDATE, DELETE, or SELECT**. A transaction must be executed **completely or not at all** to maintain database consistency. Transactions are essential to **ensure reliability and integrity** in multi-user and concurrent environments.

---

### **Properties of a Transaction (ACID Properties):**

#### **1. Atomicity:**

- Ensures that a transaction is **all-or-nothing**.
- If any part of the transaction fails, the **entire transaction is rolled back**.
- Example: Transferring money from one account to another – both debit and credit must succeed or fail together.

#### **2. Consistency:**

- Ensures that a transaction **transforms the database from one consistent state to another**.
- Integrity constraints and rules are preserved before and after the transaction.

#### **3. Isolation:**

- Ensures that **concurrent transactions** do not interfere with each other.

- Intermediate states of a transaction are **invisible to other transactions** until it commits.

#### **4. Durability:**

- Once a transaction commits, its changes are **permanent**, even in case of **system failure**.
  - DBMS ensures committed data is safely stored on **non-volatile storage**.
- 

#### **Key Points / Advantages:**

- Transactions **maintain data integrity and correctness**.
- Enable **safe concurrent access** in multi-user environments.
- Allow **recovery from failures** without loss of committed data.

(c) Write a note on view serializability (7 mark)

#### **Introduction / Definition:**

In a Database Management System (DBMS), **concurrent execution of transactions** is essential for **multi-user environments**. However, it can lead to **inconsistencies** if not managed properly. To ensure **correctness**, schedules of transactions are analyzed for **serializability**, which guarantees that the outcome of concurrent execution is the same as some **serial execution** of transactions.

**View serializability** is a **type of serializability** based on the **view of data items** accessed and modified by transactions. A schedule is **view serializable** if it is **equivalent to a serial schedule** in terms of **read and write operations on database items**.

---

#### **Definition:**

A schedule S is **view serializable** if there exists a serial schedule S' such that:

1. **Initial Reads are the Same:** If a transaction reads a data item initially, the corresponding transaction in the serial schedule also reads it initially.
  2. **Read-From Relation is Preserved:** If a transaction reads a value written by another transaction in S, it reads the same value in S'.
  3. **Final Writes are the Same:** The transaction that performs the final write of a data item in S also performs the final write in S'.
- 

### **Steps / Conditions for View Serializability:**

1. Identify **all read and write operations** of transactions.
  2. Check which transaction **writes the initial value** of each item.
  3. Check **which transaction reads values from others**.
  4. Ensure that the **final writes** match with some serial execution.
  5. If all conditions are met, the schedule is **view serializable**.
- 

### **Example:**

#### **Transactions:**

T1: R(A), W(A)

T2: R(A), W(A)

#### **Schedule S:**

R1(A), R2(A), W1(A), W2(A)

- **Read-From Relation:**

- T1 reads initial value of A, writes it.
  - T2 reads initial value of A (before T1's write), writes it.
  - **Final Write:**
    - W2(A) is the final write → matches some serial order T1 → T2.
  - **Conclusion:**
    - Schedule S is **view serializable** as it can be transformed to a serial schedule **T1 → T2** preserving initial reads, read-from relations, and final writes.
- 

### **Key Features:**

- Focuses on **view of data** rather than strict order of operations.
- Allows **more concurrency** compared to conflict serializability.
- Useful for **optimistic concurrency control** techniques.

### **Advantages:**

- Enables **safe concurrent execution** of transactions.
- Ensures **database consistency** without enforcing strict order of operations.

### **Disadvantages:**

- Harder to **check automatically** compared to conflict serializability.
  - Requires **analyzing read-from and final write relations** for all data items.
- 

### **Applications:**

- Multi-user database environments to **maximize concurrency**.
- Optimistic concurrency control methods in **distributed databases**.
- Ensuring **data consistency** without reducing transaction throughput.

**OR**

**Q.4 (a) With the help of example explain Generalization.. (3 mark)**

## **Generalization in DBMS**

### **Introduction / Definition:**

**Generalization** is a process in database design where **two or more lower-level entity sets** are combined into a **higher-level entity set** based on **common attributes**. It is the **reverse of specialization** and is used to **reduce redundancy** and create a more abstract view of data. Generalization helps in **simplifying the database structure** and representing **common characteristics** at a higher level.

---

### **Explanation:**

- **Bottom-up approach:** Start from **specific entities** → combine them into a **general entity**.
  - The general entity contains **common attributes** of all specialized entities.
  - Specialized entities may have additional attributes unique to them.
- 

### **Example:**

- **Specific Entities:**
  - Car(Vehicle\_Id, Model, Engine\_Type)

- Truck(Vehicle\_Id, Model, Load\_Capacity)
  - **Generalized Entity:**
    - Vehicle(Vehicle\_Id, Model) → captures **common attributes**
  - **Interpretation:**
    - Car and Truck are **specializations** of Vehicle.
    - Vehicle represents the **generalized entity** with attributes shared by both.
- 

### **Key Points / Advantages:**

- Reduces **redundancy** by storing **common attributes once**.
- Simplifies **database design and maintenance**.
- Helps in **hierarchical organization** of entities.

(b) Define various approach used for Log based recovery (4 mark)

### **Introduction / Definition:**

**Log-based recovery** is a method used in Database Management Systems (DBMS) to **restore the database to a consistent state** after a system crash or failure. The DBMS maintains a **log file** that records all **database modifications**. During recovery, the log is used to **redo committed transactions** and **undo uncommitted transactions**, ensuring **data integrity and durability**.

---

### **Approaches Used for Log-Based Recovery:**

1. **UNDO Logging (Rollback Logging):**
  - Records **old values** of data items before they are modified.

- If a transaction fails or aborts, the DBMS **uses the log to restore old values.**
- Ensures **atomicity** by undoing incomplete transactions.
- **Example:** If T1 updates salary from 50,000 → 55,000 and fails, use log to revert to 50,000.

## 2. REDO Logging (Forward Logging):

- Records **new values** of data items after they are modified.
- If the system crashes **after commit but before changes are written to disk**, the DBMS **reapplies changes from the log.**
- Ensures **durability** of committed transactions.
- **Example:** T2 commits updating salary to 60,000; after crash, redo ensures value is 60,000.

## 3. UNDO/REDO Logging (Combined Approach):

- Maintains both **old and new values** in the log.
- Allows the DBMS to **undo uncommitted transactions** and **redo committed transactions.**
- Provides **complete recovery** from system failures.

## 4. Checkpoints:

- Periodically writes **all committed data and log info** to disk.
- Reduces the number of log records to be scanned during recovery.
- Speeds up **recovery process** after crash.

**Key Points / Advantages:**

- Ensures **Atomicity, Consistency, and Durability** of transactions.
- Enables **fast recovery** in case of system failures.
- Minimizes **data loss and inconsistencies**.

**(c) List and explain various types of Locks used for concurrency control. (7 mark)**

### **Introduction / Definition:**

In a multi-user database environment, **concurrent transactions** may access and modify the same data simultaneously. This can lead to **inconsistencies and anomalies** such as lost updates, dirty reads, and uncommitted data. To manage this, DBMS uses **locks**, which are mechanisms that **restrict access to data items** during transaction execution. Locks ensure **concurrency control** and maintain **database consistency**.

---

### **Types of Locks:**

#### **1. Shared Lock (S-Lock / Read Lock):**

- Allows a transaction to **read a data item** but not modify it.
  - Multiple transactions can **simultaneously hold shared locks** on the same data item.
  - Prevents **write operations** until all shared locks are released.
  - **Example:** Transaction T1 reads a record; T2 can also read but cannot write until T1 releases the lock.
- 

#### **2. Exclusive Lock (X-Lock / Write Lock):**

- Allows a transaction to **both read and write a data item**.

- Only **one transaction can hold an exclusive lock** on a data item at a time.
  - Prevents **other transactions from reading or writing** the locked item until released.
  - **Example:** T1 updates a record; T2 must wait until T1 commits or rolls back.
- 

### 3. Intention Locks:

- Used in **multi-level locking** to indicate a transaction's intention to acquire **shared or exclusive locks** on lower-level items.
  - Types include:
    - **Intention Shared (IS):** Intends to acquire shared locks on some subordinate items.
    - **Intention Exclusive (IX):** Intends to acquire exclusive locks on some subordinate items.
  - Helps in **hierarchical locking** for tables, pages, and rows efficiently.
- 

### 4. Read Lock / Write Lock (Row-Level Locking):

- **Read Lock:** Allows **reading only** at row level; other reads allowed, writes blocked.
  - **Write Lock:** Allows **updating a specific row**; all other reads and writes are blocked until released.
  - Ensures **fine-grained control** and improves **concurrency**.
- 

### 5. Binary Lock:

- A **single lock bit** for a data item; either locked (1) or unlocked (0).
  - Simple but **less flexible**, can reduce concurrency.
- 

## 6. Shared-Exclusive Locks (Two-Phase Locking Protocol):

- Combines **shared and exclusive locks**.
  - Transactions **acquire locks in growing phase** and **release in shrinking phase**.
  - Ensures **conflict serializability** and prevents **lost updates**.
- 

### Key Points / Advantages:

- Locks prevent **inconsistent reads and writes** in concurrent transactions.
- Enable **conflict serializability** of schedules.
- Support **two-phase locking protocol** for database correctness.

### Disadvantages:

- May lead to **deadlocks** if multiple transactions wait indefinitely.
- Requires **careful management** to avoid performance bottlenecks.

**Q.5 (a)** How can we print output in PL/SQL block? Give example. (3 mark)

### Introduction / Definition:

In PL/SQL, a **block** is a unit of code that can include declarations, executable statements, and exception handling. To **display output** from a PL/SQL block, the **DBMS\_OUTPUT.PUT\_LINE** procedure

is commonly used. It allows printing messages or variable values to the console or output window.

---

### Syntax:

DBMS\_OUTPUT.PUT\_LINE('message or value');

- 'message or value' → The text or variable value to be printed.
  - Multiple calls to PUT\_LINE print **each line separately**.
- 

### Example:

BEGIN

```
DBMS_OUTPUT.PUT_LINE('Hello, GTU Students!');

DBMS_OUTPUT.PUT_LINE('Welcome to PL/SQL
programming.');

END;
```

### Explanation:

- The BEGIN...END block contains **executable statements**.
- Each DBMS\_OUTPUT.PUT\_LINE prints a line to the output.
- Output:

Hello, GTU Students!

Welcome to PL/SQL programming.

---

### Key Points / Advantages:

- Useful for **debugging and testing** PL/SQL code.

- Can print **variables, expressions, or messages.**
- Multiple lines can be printed sequentially.

**(b) Explain cursor and its types. (4 mark)**

### **Introduction / Definition:**

A **cursor** in PL/SQL is a **pointer that allows row-by-row processing of SQL query results**. It is used to **retrieve, manipulate, and process individual rows** from the result set of a SELECT statement. Cursors are essential when **multiple rows are returned**, and we need **control over each row**.

---

### **Types of Cursors:**

#### **1. Implicit Cursor:**

- Automatically created by PL/SQL when a **SELECT INTO, INSERT, UPDATE, or DELETE** statement is executed.
- Suitable for **single-row queries**.
- PL/SQL automatically manages the cursor lifecycle (open, fetch, close).
- **Example:**

DECLARE

v\_salary NUMBER;

BEGIN

    SELECT salary INTO v\_salary FROM Employee WHERE eid =  
    101;

    DBMS\_OUTPUT.PUT\_LINE('Salary: ' || v\_salary);

END;

/

- Here, the cursor is **implicit**, handled by PL/SQL internally.
- 

## 2. Explicit Cursor:

- Defined by the programmer for **multi-row queries**.
- Requires **manual operations**: OPEN, FETCH, and CLOSE.
- Allows **row-by-row processing**.
- **Example:**

```
DECLARE
```

```
    CURSOR emp_cursor IS
        SELECT ename, salary FROM Employee;
        v_name Employee.ename%TYPE;
        v_salary Employee.salary%TYPE;

    BEGIN
        OPEN emp_cursor;
        LOOP
            FETCH emp_cursor INTO v_name, v_salary;
            EXIT WHEN emp_cursor%NOTFOUND;
            DBMS_OUTPUT.PUT_LINE('Employee: ' || v_name || ', Salary: '
                || v_salary);
        END LOOP;
        CLOSE emp_cursor;
    END;
```

---

```
/
```

## **Key Points / Advantages:**

- Cursors allow **processing one row at a time** in PL/SQL.
  - Implicit cursors are **simple for single-row queries**.
  - Explicit cursors are **flexible for multi-row queries**.
  - Cursors provide **control over query execution** and enable **conditional processing**.
- 

(c) Write SQL queries for the following (query 1,2 and 3 carry one mark/query 4 & 5 carry two marks) Employee(eid, ename, salary, deptid) Department(deptid, dname)

1. List name of all employee
2. List all the departments with id > 5
3. List sum of the salary all employees.
4. Department wise count the number of employees
5. Find the name of employee earning highest salary.

## **SQL Queries for Employee and Department Tables**

### **Given Tables:**

- Employee(eid, ename, salary, deptid)
  - Department(deptid, dname)
- 

### **1. List name of all employees (1 Mark)**

SELECT ename

FROM Employee;

### **Explanation:**

- Uses SELECT to fetch the ename column from the Employee table.
- 

## **2. List all the departments with id > 5 (1 Mark)**

```
SELECT *  
FROM Department  
WHERE deptid > 5;
```

### **Explanation:**

- Fetches all columns from Department where deptid is greater than 5.
- 

## **3. List sum of the salary of all employees (1 Mark)**

```
SELECT SUM(salary) AS total_salary  
FROM Employee;
```

### **Explanation:**

- Uses the aggregate function SUM to calculate **total salary** of all employees.
- 

## **4. Department-wise count of the number of employees (2 Marks)**

```
SELECT deptid, COUNT(*) AS emp_count  
FROM Employee  
GROUP BY deptid;
```

### **Explanation:**

- Groups employees by deptid and counts the number of employees in each department.

- GROUP BY is used for **aggregate calculations per department**.
- 

## 5. Find the name of employee earning the highest salary (2 Marks)

```
SELECT ename  
FROM Employee  
WHERE salary = (SELECT MAX(salary) FROM Employee);
```

### Explanation:

- Subquery (SELECT MAX(salary) FROM Employee) finds the **highest salary**.
  - Outer query fetches the employee(s) with that salary.
  - Ensures **correct result even if multiple employees share the highest salary**.
- 

### Key Points:

- SQL aggregate functions like SUM, COUNT, and MAX are used for **calculations**.
- GROUP BY helps in **department-wise aggregation**.
- Subqueries allow **filtering based on computed values**.

**OR**

**Q.5 (a) Define encryption and decryption. (3 mark)**

### Introduction

/

### Definition:

**Encryption** and **decryption** are fundamental concepts in **cryptography**, which is used to secure data from unauthorized access. These techniques transform data between **readable (plaintext)** and

**unreadable (ciphertext)** forms to ensure **confidentiality** and **security**.

---

### **Encryption:**

- Encryption is the process of **converting plaintext (original data)** into **ciphertext (encoded data)** using an **encryption algorithm and key**.
  - Purpose: To **protect sensitive information** during storage or transmission.
  - **Example:**
    - Plaintext: "HELLO"
    - Ciphertext (using a simple shift key 3): "KHOOR"
- 

### **Decryption:**

- Decryption is the **reverse process of encryption**, converting **ciphertext back to plaintext** using a **decryption algorithm and key**.
  - Ensures that only **authorized users** can access the original data.
  - **Example:**
    - Ciphertext: "KHOOR"
    - Plaintext (after decryption): "HELLO"
- 

### **Key Points / Advantages:**

- Ensures **data confidentiality** during storage and transmission.
- Forms the basis for **secure communication, digital signatures, and authentication**.

- Works with **symmetric keys** (same key for encryption and decryption) or **asymmetric keys** (public and private keys).

(b) List various types of Triggers. (4 mark)

### **Introduction**

/

### **Definition:**

A **trigger** is a **PL/SQL block** that automatically executes in response to a **specific event** on a table or view. Triggers are used to **enforce business rules, maintain integrity, or audit changes** in the database. They execute **automatically** without explicit invocation.

---

### **Types of Triggers:**

#### **1. BEFORE Trigger:**

- Executes **before the triggering event** occurs (INSERT, UPDATE, DELETE).
- Often used to **validate or modify data** before it is written to the table.
- **Example:** Check if salary is above a minimum before inserting a record.

#### **2. AFTER Trigger:**

- Executes **after the triggering event** has occurred.
- Commonly used for **auditing, logging, or updating related tables**.
- **Example:** Log changes in a separate audit table after an update.

#### **3. INSTEAD OF Trigger:**

- Used mainly with **views**.
- Executes **instead of the triggering action** to allow **modification of views** that cannot be updated directly.

- **Example:** Update underlying table columns through a view.

#### **4. ROW-Level Trigger:**

- Executes **once for each row affected** by the triggering event.
- Useful for **row-specific validation or processing**.
- **Example:** Automatically calculate total salary per employee row during insert.

#### **5. STATEMENT-Level Trigger:**

- Executes **once per SQL statement**, regardless of the number of rows affected.
  - Efficient for operations that are **not row-specific**, like logging the number of rows affected.
- 

#### **Key Points / Advantages:**

- Enforces **business rules automatically**.
- Maintains **data integrity and auditing**.
- Reduces the need for **manual intervention** in repetitive tasks.

(c) What is role? How can we authorize a user using grant and revoke command? Give example. (7 mark)

#### **Introduction**

/

#### **Definition:**

In a Database Management System (DBMS), a **role** is a **named collection of privileges** that can be assigned to users. Roles simplify **user management and access control** by grouping related privileges together. Instead of granting individual privileges to each user, a role can be **granted once and assigned to multiple users**, making administration more efficient.

---

## Granting Privileges Using Roles

- **GRANT** command is used to **assign privileges or roles** to users or other roles.
- Privileges can include **SELECT, INSERT, UPDATE, DELETE** on tables, or **system privileges** like CREATE TABLE.
- **Syntax:**

GRANT privilege\_or\_role TO user\_name [WITH ADMIN OPTION];

### Example:

1. Create a role:

CREATE ROLE manager\_role;

2. Grant privileges to role:

GRANT SELECT, INSERT, UPDATE ON Employee TO manager\_role;

3. Assign role to a user:

GRANT manager\_role TO user1;

- **Explanation:** User user1 can now **select, insert, and update** Employee table because of the role.

---

## Revoking Privileges Using Roles

- **REVOKE** command is used to **remove privileges or roles** from users or other roles.
- **Syntax:**

REVOKE privilege\_or\_role FROM user\_name;

### Example:

REVOKE manager\_role FROM user1;

- **Explanation:** Removes all privileges associated with manager\_role from user1.
  - Useful to **control access** when users change roles or responsibilities.
- 

### **Key Points / Advantages:**

- **Simplifies privilege management:** Assigning roles is easier than granting individual privileges.
- **Improves security:** Privileges can be **controlled centrally** using roles.
- **Flexible:** Roles can be **granted to multiple users** and **revoke privileges easily**.